EA-MT Entity-Aware Machine Translation in English to Chinese

Monday 25th November, 2024

Jou-Yi Lee
jlee1039@ucsc.edu

Abstract

The primary objective of this project is to tackle the challenge of translating a given input sentence from the source language, English, into the target language, Chinese. This task focuses on addressing the accurate translation of named entities, which may be rare, ambiguous, or unknown to traditional machine translation systems.

The goal is to develop a machine translation system capable of effectively handling such named entities within input sentences, ensuring that their translation into the target language is precise.

1 Introduction

Named entity translation presents a unique set of challenges in machine translation systems due to their reliance on contextual, structural, and linguistic nuances. Ambiguity in context is a pervasive issue, as entities often depend on surrounding text for proper interpretation. For instance, the term "*Apple*" may refer to the fruit or the technology company, and distinguishing between these interpretations becomes particularly challenging when the contextual clues are limited or absent.

Another complexity arises from multi-word entities, such as names of organizations, products, or events, which require precise handling to maintain semantic accuracy. Literal translations of such entities often lead to awkward or incorrect results. For example, translating "*Department of Justice*" word-for-word fails to convey the appropriate equivalent term in the target language, thus diminishing the overall quality of the translation.

Cross-lingual consistency further complicates the translation process, as entities may have multiple valid representations in the target language. Maintaining consistency across translations is critical to ensuring that the intended meaning is preserved in various linguistic contexts.

To address these challenges, this study employs the **IWSLT 2017 English-to-Chinese** dataset, focusing on named entity translation. Named entities in the dataset are first identified using **SpaCy**, a robust named entity recognition (NER) tool. These recognized entities are then incorporated into a machine translation pipeline based on the **mBERT** architecture, leveraging pre-trained and fine-tuned weights to enhance the translation accuracy of complex and ambiguous entities. This methodology seeks to establish a robust framework for improving named entity translation, ensuring semantic and contextual fidelity in cross-lingual tasks.

2 IWSLT Dataset

The International Conference on Spoken Language Translation (IWSLT) 2017 dataset involves English-to-Chinese translations and provides a valuable resource for real-world conversational examples. For instance:

> *"Thank you so much, Chris. And it's truly a great honor to have the opportunity to come to this stage twice; I'm extremely grateful."*
>
> "非常谢谢，克里斯。的确非常荣幸能有第二次站在这个台上的机会，我真是非常感激。"

In this dataset, English sentence lengths are more dispersed compared to Chinese sentences. Chinese sentence lengths are concentrated

in shorter intervals because Chinese is a character-based language and typically uses fewer words.

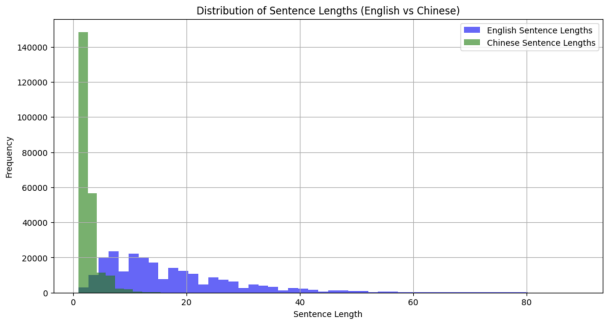

图 1: Distribution of Sentence Lengths (English vs. Chinese)

2.1 Linguistic Structural Differences

- **Character-Based vs. Word-Based Languages:**

  Chinese, as a character-based language, treats each character as the smallest semantic unit. In contrast, English, as a word-based language, requires additional auxiliary words (e.g., articles, prepositions, auxiliary verbs) to form complete sentences. Common examples of auxiliary words in English include articles (*"the"*), prepositions (*"in"*), and auxiliary verbs (*"is"*).

  For instance:

  – English: *"The dog is on the mat."* (6 words)

  – Chinese: "狗在墊子上。" (5 characters)

  This demonstrates that Chinese often uses fewer units to convey the same meaning, while English relies on additional grammatical components to structure sentences.

- **Differences in Word Order and Syntax:**

  English generally follows a fixed Subject-Verb-Object (SVO) structure, which makes sentence construction predictable and uniform. On the other hand, Chinese exhibits more flexibility in word order, allowing sentence structure to vary depending on the intended emphasis. However, this flexibility in Chinese may lead to alignment challenges during translation, especially when attempting to map flexible Chinese structures to the fixed SVO patterns of English.

  For example:

  – English: *"The red book on the table belongs to me."*

  – Chinese: "桌上的紅色書是我的。" (Reordered emphasis on *"the red book on the table"*)

  Furthermore, Chinese allows for reordering of elements to emphasize different parts of a sentence:

  – Standard: "我喜歡這本書。" (Subject-first: *"I like this book."*)

  – Emphasis: "這本書我喜歡。" (Emphasizing *"this book"*)

These structural differences pose challenges for machine translation models. While English requires adherence to strict grammatical structures, Chinese translation must account for its compact nature and flexible word order. Misalignment of sentence structures during translation can result in semantic or syntactic inaccuracies. Models must, therefore, be designed to accommodate the unique linguistic characteristics of both languages to achieve high-quality translations.

IWSLT English Word Cloud of Tokens in the Dataset

The word cloud illustrates the most common tokens within the IWSLT English-to-Chinese translation dataset. As expected, function words such as "the," "and," "to," "of," and "a" are prominently featured due to their frequent usage in the English language. This observation aligns with Zipf's Law, which states that in most natural language texts, function words tend to appear far more frequently than semantic words.

Furthermore, the word cloud highlights tokens such as "can," "they," and "this," which suggest the presence of conversational or descriptive text within the dataset. Although function words dominate the visualization, less frequent tokens like "what," "years," and "think" indicate potential narrative or explanatory content.

图 2: Word Cloud of Tokens in the IWSLT Dataset

IWSLT Chinese Word Cloud of Tokens in the Dataset

As shown in the visualization, high-frequency tokens such as “的,” “是,” “了,” “在,” and “有” dominate the dataset. These words are function words in Chinese and play a key role in sentence structure, similar to their English counterparts like “the,” “and,” or “to.” This observation aligns with Zipf’s Law, which posits that functional words frequently outnumber semantic words in natural language texts.

The visualization also highlights additional words like “我们,” “他们,” and “自己,” which indicate conversational or narrative contexts within the dataset. While function words are predominant, other tokens such as “因为,” “可以,” and “为什么” suggest the presence of explanatory or argumentative content.

图 3: Word Cloud of Tokens in the IWSLT Dataset (Chinese)

The analysis reveals that both character-based and word-based languages follow Zipf’s Law. This law highlights that in natural language, a small number of high-frequency tokens dominate the dataset, while the majority of tokens appear infrequently. The word cloud further reflects the diversity of the IWSLT dataset, suggesting that it likely contains a mix of narrative, conversational, and explanatory text types. The analysis reveals that both character-based and word-based languages follow Zipf’s Law. This law highlights that in natural language, a small number of high-frequency tokens dominate the dataset, while the majority of tokens appear infrequently. This visualization indicates that the dataset includes diverse text genres, such as narrative, conversational, and explanatory elements. It offers a preliminary understanding of the dataset’s composition and variety. It can guide preprocessing strategies, such as removing stopwords (e.g., “The", "a",“的,” “是,” “了,”) to focus on more semantically meaningful tokens for downstream natural language processing tasks.

3 Methodology

3.1 Flowchart

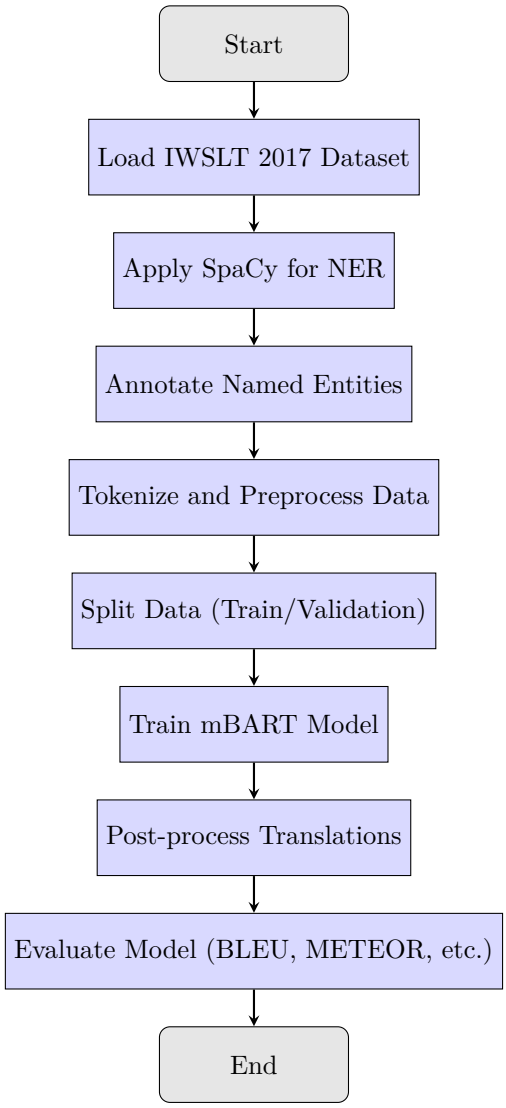

3.2 Data Preprocessing

In this study, the **IWSLT** 2017 English-to-Chinese translation dataset was utilized, with a focus on addressing the challenge of translating named entities. To enhance the model's ability to process named entities, additional steps were taken to improve entity recognition. Specifically, **SpaCy's** Transformer-based English named entity recognition model (`en_core_web_trf`) was employed to identify entities within the source sentences. The identified entities were annotated in the format «`{ent.label}:{ent.text}>` , for example, «`ORG:Apple`». This annotation allows the model to recognize these key segments during training and inference. This preprocessing step significantly enhances the model's awareness of named entities, thereby optimizing translation quality.

3.3 SpaCy

未完成 1. **識Ⓕ完成的資料** 2. 資料增Ⓕ (Data Augmentation)：數據集進行了增Ⓕ處理，例如隨機替Ⓕ同義詞、分詞重組、或添加噪聲。3. Ⓕ明添加實體標記對於模型上下文學習的幫助，以及如何處理不同語言中實體對應的差Ⓕ性。4. 不平衡數據處理：如果數據集中某些類型句子或語言分Ⓕ不均，Ⓕ明是否Ⓕ用了重新加權或欠Ⓕ樣/過Ⓕ樣策略。

3.4 Model Design — mBART

The translation system was built upon the mBART model (`mbart-large-50-many-to-many-mmt`), a pre-trained Transformer-based sequence-to-sequence model designed for multilingual tasks. The tokenizer of the mBART model was customized to correctly recognize and process inputs with entity markers. By fine-tuning the preprocessed dataset, the model was able to leverage the additional contextual information provided by the named entity annotations, thereby improving translation performance.

未完成 **Ⓕ明模型** 1. 詳細Ⓕ明 mBART 的架構（Transformer-based），及其對多語言任務的適應性。2. 明預訓練權重的重要性及其在遷移學習中的優勢。3. Ⓕ**明參數**提供關鍵參數，例如學習率、批量大小、早停條件和訓練次數。解釋Ⓕ什Ⓕ使用交叉熵損失和標Ⓕ平滑的原因，ⒻⒻ調其對模型穩定性的貢獻。4. Ⓕ**明權重**

5. **注意力分Ⓕ (Attention Visualization)：**用於分析模型如何處理語言結構，特Ⓕ是命名實體。6. Ⓕ**明訓練過程**訓練環境: GPU 與訓練時間 7. Ⓕ**明 Loss**

8. **Precision Recall F1** 9. **模型比較** 1. Ⓕ有訓練過的 mBART 與訓練過的比較改進幅度, 用不同的評估指標。2. mBART 與其他翻譯模型（如 Google NMT、Transformer-based 模型）的優劣勢。

4 Evaluation Metrics

4.1 METEOR

We chose METEOR as our evaluation metric, which assesses machine translation quality through unigram matching with human references. METEOR compares unigrams by surface form, stem, and meaning, calculating scores based on precision, recall, and fragmentation. METEOR scores translations based on unigram precision, unigram recall, and fragmentation, the latter of which captures the word order accuracy in relation to the reference.

4.2 BLUE

We will also use BLEU for contrast. BLEU evaluates machine-translated text by comparing it to human translations, with higher scores indicating closer alignment to professional human translation. One of the first metrics to correlate highly with human judgment, BLEU remains popular and cost-effective. It calculates scores for individual segments, typically sentences, by comparing them to reference translations, then averages these scores across the corpus for an overall quality estimate. However, BLEU does not account for intelligibility or grammatical correctness.

5 Proposed Approach and Methodology

The process of integrating entity-aware translation is divided into several steps to manage the complexities of accurately translating named entities from English to Chinese.